

DSC680-Project1

September 27, 2025

```
[1]: # 1) Setup
import os, zipfile
from pathlib import Path
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.linear_model import LinearRegression, RidgeCV
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
import joblib

import warnings
warnings.filterwarnings("ignore")
from IPython.core.display import display, HTML
display(HTML("<style>.jp-OutputArea { overflow: visible !important; }</style>"))

dataset_path = r"C:
    ↵\Users\amyha\OneDrive\Documents\DSC680\car_insurance_premium_dataset.csv"
df = pd.read_csv(dataset_path)

<IPython.core.display.HTML object>
```


```
[2]: # 2) Summary statistics
display(df.describe(include='all'))
```

| | Driver Age | Driver Experience | Previous Accidents | \ |
|-------|-------------|-------------------|--------------------|---|
| count | 1000.000000 | 1000.000000 | 1000.0000 | |
| mean | 41.575000 | 14.759000 | 2.5680 | |
| std | 13.765677 | 10.544292 | 1.6989 | |
| min | 18.000000 | 0.000000 | 0.0000 | |
| 25% | 30.000000 | 6.000000 | 1.0000 | |
| 50% | 42.000000 | 13.000000 | 3.0000 | |
| 75% | 53.000000 | 23.000000 | 4.0000 | |


| | | | | |
|-------|---------------------------|------------------------|-------------|---|
| max | 65.000000 | 40.000000 | 5.0000 | |
| | Annual Mileage (x1000 km) | Car Manufacturing Year | Car Age | \ |
| count | 1000.000000 | 1000.000000 | 1000.000000 | |
| mean | 17.933000 | 2007.637000 | 17.363000 | |
| std | 4.410665 | 10.363331 | 10.363331 | |
| min | 11.000000 | 1990.000000 | 0.000000 | |
| 25% | 14.000000 | 1999.000000 | 8.000000 | |
| 50% | 18.000000 | 2008.000000 | 17.000000 | |
| 75% | 22.000000 | 2017.000000 | 26.000000 | |
| max | 25.000000 | 2025.000000 | 35.000000 | |
| | Insurance Premium (\$) | | | |
| count | 1000.000000 | | | |
| mean | 493.742250 | | | |
| std | 5.909689 | | | |
| min | 477.050000 | | | |
| 25% | 489.487500 | | | |
| 50% | 493.950000 | | | |
| 75% | 498.312500 | | | |
| max | 508.150000 | | | |

```
[3]: # 3) Distributions - one chart per column
numeric_cols = [c for c in df.columns if pd.api.types.is_numeric_dtype(df[c])]
for col in numeric_cols:
    plt.figure()
    df[col].plot(kind='hist', bins=20, title=f'Histogram: {col}')
    plt.xlabel(col)
    plt.ylabel('Count')
    plt.show()
```


Histogram: Driver Age




Histogram: Driver Experience




Histogram: Previous Accidents







Histogram: Car Manufacturing Year




Histogram: Car Age






```
[4]: # 4) Box plots - one chart per column
for col in numeric_cols:
    plt.figure()
    df.boxplot(column=col)
    plt.title(f'Boxplot: {col}')
    plt.ylabel(col)
    plt.show()
```


Boxplot: Driver Age




Boxplot: Driver Experience




Boxplot: Previous Accidents




Boxplot: Annual Mileage (x1000 km)




Boxplot: Car Manufacturing Year




Boxplot: Car Age





```
[5]: # 5) Correlation heatmap
corr = df[numeric_cols].corr(numeric_only=True)
plt.figure(figsize=(6,5))
im = plt.imshow(corr.values, interpolation='nearest')
plt.title('Correlation Heatmap')
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.index)), corr.index)
plt.colorbar(im, fraction=0.046, pad=0.04)
plt.tight_layout()
plt.show()
display(corr)
```




| | Driver Age | Driver Experience | Previous Accidents | \ |
|---------------------------|---------------------------|------------------------|--------------------|---|
| Driver Age | 1.000000 | 0.607890 | 0.031819 | |
| Driver Experience | 0.607890 | 1.000000 | 0.020837 | |
| Previous Accidents | 0.031819 | 0.020837 | 1.000000 | |
| Annual Mileage (x1000 km) | 0.056822 | -0.014424 | 0.007088 | |
| Car Manufacturing Year | 0.008187 | -0.038194 | -0.030123 | |
| Car Age | -0.008187 | 0.038194 | 0.030123 | |
| Insurance Premium (\$) | -0.776848 | -0.803323 | -0.030123 | |
| | Annual Mileage (x1000 km) | Car Manufacturing Year | \ | |
| Driver Age | 0.056822 | 0.008187 | | |
| Driver Experience | -0.014424 | -0.038194 | | |
| Previous Accidents | 0.007088 | -0.030123 | | |
| Annual Mileage (x1000 km) | 1.000000 | -0.002898 | | |
| Car Manufacturing Year | -0.002898 | 1.000000 | | |
| Car Age | 0.002898 | -1.000000 | | |
| Insurance Premium (\$) | 0.022131 | -0.171829 | | |

| | Car Age | Insurance Premium (\$) |
|---------------------------|-----------|------------------------|
| Driver Age | -0.008187 | -0.776848 |
| Driver Experience | 0.038194 | -0.803323 |
| Previous Accidents | 0.030123 | 0.410786 |
| Annual Mileage (x1000 km) | 0.002898 | 0.022131 |
| Car Manufacturing Year | -1.000000 | -0.171829 |
| Car Age | 1.000000 | 0.171829 |
| Insurance Premium (\$) | 0.171829 | 1.000000 |

```
[6]: # 6) Scatter Plot
plt.figure(figsize=(6,4))
sns.scatterplot(x=df["Driver Experience"],
                 y=df["Insurance Premium ($)"],
                 alpha=0.5, color="orange")

sns.regplot(x=df["Driver Experience"],
            y=df["Insurance Premium ($)"],
            scatter=False, color="red")

plt.title("Premium vs Driver Experience")
plt.xlabel("Experience (years)")
plt.ylabel("Premium ($)")
plt.show()
```



```
[12]: # 7) Build and test models
TARGET = "Insurance Premium ($)"
assert TARGET in df.columns, f"Target column '{TARGET}' not found."
X = df.drop(columns=[TARGET])
y = df[TARGET]

num_cols = X.select_dtypes(include=["number"]).columns.tolist()
cat_cols = X.select_dtypes(exclude=["number"]).columns.tolist()

numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler(with_mean=False))
])

categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preprocessor = ColumnTransformer(transformers=[
    ("num", numeric_transformer, num_cols),
    ("cat", categorical_transformer, cat_cols)
])

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42
)

models = {
    "LinearRegression": LinearRegression(),
    "RidgeCV": RidgeCV(alphas=np.logspace(-2, 2, 9)), # 0.01..100
    "RandomForest": RandomForestRegressor(n_estimators=200, random_state=42),
    "GradientBoosting": GradientBoostingRegressor(n_estimators=250, learning_rate=0.1, max_depth=3, random_state=42),
}

results = []
best_model = None
best_name = None
best_rmse = float("inf")

for name, model in models.items():
    pipe = Pipeline(steps=[("prep", preprocessor), ("model", model)])
    pipe.fit(X_train, y_train)
    y_pred = pipe.predict(X_test)
    rmse = mean_squared_error(y_test, y_pred, squared=False)
    mae = mean_absolute_error(y_test, y_pred)
```

```

r2 = r2_score(y_test, y_pred)
results.append({"model": name, "test_rmse": rmse, "test_mae": mae, "test_r2": r2})
if rmse < best_rmse:
    best_rmse = rmse
    best_model = pipe
    best_name = name

results_df = pd.DataFrame(results).sort_values("test_rmse").
    reset_index(drop=True)
display(results_df)

```

| | model | test_rmse | test_mae | test_r2 |
|---|------------------|--------------|--------------|----------|
| 0 | LinearRegression | 5.084230e-14 | 3.808509e-14 | 1.000000 |
| 1 | RidgeCV | 5.499015e-05 | 4.545626e-05 | 1.000000 |
| 2 | GradientBoosting | 3.598881e-01 | 2.782540e-01 | 0.996182 |
| 3 | RandomForest | 7.357459e-01 | 5.486713e-01 | 0.984044 |


[13]: # 8) Residual Plot

```

y_pred_best = best_model.predict(X_test)
residuals = y_test - y_pred_best

plt.figure(figsize=(6,4))
plt.scatter(y_pred_best, residuals, alpha=0.6)
plt.axhline(0, linestyle="--")
plt.title(f"Residuals vs Predicted ({best_name})")
plt.xlabel("Predicted Premium")
plt.ylabel("Residual")
plt.tight_layout()
plt.show()

```



```
[15]: # 9) Save best model and example inference
model_path = f'best_model_{best_name.replace(" ", "_").lower()}.joblib'
joblib.dump(best_model, model_path)
print('Best model:', best_name)

example = pd.DataFrame([
    col: (X_train[col].median() if pd.api.types.is_numeric_dtype(X_train[col])_
    ↵else X_train[col].mode().iloc[0])
    for col in X_train.columns
])
display(example)
print('Prediction for example row:')
print(best_model.predict(example))
```

Best model: LinearRegression

| | Driver Age | Driver Experience | Previous Accidents | \ |
|---|------------|-------------------|--------------------|---|
| 0 | 42.0 | 13.0 | 3.0 | |

| | Annual Mileage (x1000 km) | Car Manufacturing Year | Car Age |
|---|---------------------------|------------------------|---------|
| 0 | 18.0 | 2008.0 | 17.0 |

Prediction for example row:
[494.8]

[]: