

# DSC680-Project3

November 22, 2025

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import zipfile
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
import numpy as np

import warnings
warnings.filterwarnings("ignore")
from IPython.core.display import display, HTML
display(HTML("<style>.jp-OutputArea { overflow: visible !important; }</style>"))

<IPython.core.display.HTML object>
```

```
[3]: # 1. Load the dataset
dataset_path = r"C:
    ↵\Users\amyha\OneDrive\Documents\DSC680\Mental_Health_and_Social_Media_Balance_Dataset.
    ↵csv"
df = pd.read_csv(dataset_path)
```

```
[4]: # 2. Exploratory Data Analysis (EDA)
# Descriptive statistics
print(df.head())
print(df.describe())
print(df.isnull().sum())
```

User_ID	Age	Gender	Daily_Screen_Time(hrs)	Sleep_Quality(1-10)	\
0	U001	44	Male	3.1	7.0
1	U002	30	Other	5.1	7.0
2	U003	23	Other	7.4	6.0
3	U004	36	Female	5.7	7.0
4	U005	34	Female	7.0	4.0

	Stress_Level(1-10)	Days_Without_Social_Media	Exercise_Frequency(week)	\
0	6.0	2.0	5.0	
1	8.0	5.0	3.0	
2	7.0	1.0	3.0	
3	8.0	1.0	1.0	
4	7.0	5.0	1.0	

	Social_Media_Platform	Happiness_Index(1-10)	
0	Facebook	10.0	
1	LinkedIn	10.0	
2	YouTube	6.0	
3	TikTok	8.0	
4	X (Twitter)	8.0	

	Age	Daily_Screen_Time(hrs)	Sleep_Quality(1-10)	\
count	500.000000	500.000000	500.000000	
mean	32.988000	5.530000	6.304000	
std	9.960637	1.734877	1.529792	
min	16.000000	1.000000	2.000000	
25%	24.000000	4.300000	5.000000	
50%	34.000000	5.600000	6.000000	
75%	41.000000	6.700000	7.000000	
max	49.000000	10.800000	10.000000	

	Stress_Level(1-10)	Days_Without_Social_Media	\
count	500.000000	500.000000	
mean	6.618000	3.134000	
std	1.542996	1.858751	
min	2.000000	0.000000	
25%	6.000000	2.000000	
50%	7.000000	3.000000	
75%	8.000000	5.000000	
max	10.000000	9.000000	

	Exercise_Frequency(week)	Happiness_Index(1-10)	
count	500.000000	500.000000	
mean	2.448000	8.376000	
std	1.428067	1.524228	
min	0.000000	4.000000	
25%	1.000000	7.000000	
50%	2.000000	9.000000	
75%	3.000000	10.000000	
max	7.000000	10.000000	


  

User_ID	0	
Age	0	
Gender	0	
Daily_Screen_Time(hrs)	0	
Sleep_Quality(1-10)	0	
Stress_Level(1-10)	0	


```
Days_Without_Social_Media      0
Exercise_Frequency(week)       0
Social_Media_Platform         0
Happiness_Index(1-10)          0
dtype: int64
```

```
[5]: # Histograms for numeric features
num_cols = [
    "Age",
    "Daily_Screen_Time(hrs)",
    "Sleep_Quality(1-10)",
    "Stress_Level(1-10)",
    "Days_Without_Social_Media",
    "Exercise_Frequency(week)",
    "Happiness_Index(1-10)"
]

plt.figure(figsize=(15, 12))
for i, col in enumerate(num_cols):
    plt.subplot(4, 2, i + 1)
    plt.hist(df[col], bins=20, edgecolor='black')
    plt.title(f"Distribution of {col}")
    plt.xlabel(col)
    plt.ylabel("Frequency")
plt.tight_layout()
plt.show()
```




```
[6]: # Correlation heatmap (encode categorical variables for correlation)
plt.figure(figsize=(10, 8))
corr = df[num_cols].corr()
sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f", square=True)
plt.title("Correlation Matrix of Numerical Features")
plt.show()
```



```
[7]: # Box plots
plt.figure(figsize=(14, 6))
plt.subplot(1, 2, 1)
sns.boxplot(x="Gender", y="Stress_Level(1-10)", data=df)
plt.title("Stress Level by Gender")

plt.subplot(1, 2, 2)
sns.boxplot(x="Social_Media_Platform", y="Stress_Level(1-10)", data=df)
plt.xticks(rotation=45)
plt.title("Stress Level by Social Media Platform")
plt.tight_layout()
plt.show()
```



```
[8]: X = df.drop(columns=["User_ID", "Stress_Level(1-10)"])
y = df["Stress_Level(1-10)"]

categorical_cols = ["Gender", "Social_Media_Platform"]
numerical_cols = X.drop(columns=categorical_cols).columns.tolist()

preprocessor = ColumnTransformer(transformers=[
    ("num", "passthrough", numerical_cols),
    ("cat", OneHotEncoder(drop="first"), categorical_cols)
])

models = {
    "Linear Regression": LinearRegression(),
    "Random Forest": RandomForestRegressor(random_state=42),
    "Gradient Boosting": GradientBoostingRegressor(random_state=42)
}

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

results = {}
for name, model in models.items():
    pipeline = Pipeline(steps=[
        ("preprocessor", preprocessor),
        ("model", model)
    ])
    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)
    results[name] = {
        "MAE": mean_absolute_error(y_test, y_pred),
        "R2": r2_score(y_test, y_pred)
    }

```

```
        "RMSE": mean_squared_error(y_test, y_pred, squared=False),  
        "R2": r2_score(y_test, y_pred)  
    }  
  
results_df = pd.DataFrame(results).T.sort_values("RMSE")  
print(results_df)
```

	MAE	RMSE	R2
Random Forest	0.797800	0.941227	0.632159
Gradient Boosting	0.814747	0.955392	0.621004
Linear Regression	0.810950	0.966219	0.612365

[ ]: